

Forward-Looking Statements



This presentation may contain forward-looking statements regarding future events, plans or the expected financial performance of our company, including our expectations regarding our products, technology, strategy, customers, markets, acquisitions and investments. These statements reflect management's current expectations, estimates and assumptions based on the information currently available to us. These forward-looking statements are not guarantees of future performance and involve significant risks, uncertainties and other factors that may cause our actual results, performance or achievements to be materially different from results, performance or achievements expressed or implied by the forward-looking statements contained in this presentation.

For additional information about factors that could cause actual results to differ materially from those described in the forward-looking statements made in this presentation, please refer to our periodic reports and other filings with the SEC, including the risk factors identified in our most recent quarterly reports on Form 10-Q and annual reports on Form 10-K, copies of which may be obtained by visiting the Splunk Investor Relations website at www.investors.splunk.com or the SEC's website at www.sec.gov. The forward-looking statements made in this presentation are made as of the time and date of this presentation. If reviewed after the initial presentation, even if made available by us, on our website or otherwise, it may not contain current or accurate information. We disclaim any obligation to update or revise any forward-looking statement based on new information, future events or otherwise, except as required by applicable law.

In addition, any information about our roadmap outlines our general product direction and is subject to change at any time without notice. It is for informational purposes only and shall not be incorporated into any contract or other commitment. We undertake no obligation either to develop the features or functionalities described or to include any such feature or functionality in a future release.

Splunk, Splunk>, Data-to-Everything, D2E and Turn Data Into Doing are trademarks and registered trademarks of Splunk Inc. in the United States and other countries. All other brand names, product names or trademarks belong to their respective owners. © 2021 Splunk Inc. All rights reserved.

Effective & Affordable Cyber-Security Vulnerability Management With Splunk Enterprise

SEC1075A

Daniel Ferreira

Vulnerability Management Lead
Global Cyber-Security Team | Shell IT International B.V.

splunk> .conf21



What if you can create your own 'natural language' search engine?



```
1 `asset` conf.splunk.com  
2 | table asset_owner asset_business_unit asset_exposure asset_exposed_ports
```

```
1 `asset` (AWS OR Azure) virtual machine internet facing  
2 | table asset_name asset_ad_domain asset_ad_memberof asset_installed_agent
```

```
1 `asset` tcp (80 OR 443 OR 8443) internet facing (BigIP OR Big-IP OR TMUI)  
2 | table a*name a*type a*owner a*country* a*ports a*last_seen a*service
```

```
1 `asset` "solarwinds agent"  
2 | table a*name a*type a*owner a*platform a*last_logged_on
```

```
1 `software` apache struts 2.* internet facing  
2 | table asset_name a*exposure software_vendor software_name s*_version
```

```
1 `identity` (domain OR enterprise) admins  
2 | table identity_forest i*_domain i*_name i*_owner i*_memberof i*_created i*logon*
```

```
1 `mobile` noncompliant (iPad OR iPhone) 14.*  
2 | table mobile_imei mobile_platform mobile_model mobile_os mobile_owner mobile_last_seen
```

```
1 `vulnerability` CVE-2021-34527 PrintNightmare (remediated OR open)  
2 | table asset_name vulnerability_unique_id v*remediation_status v*first_seen v*last_seen
```



Daniel Ferreira

Vulnerability Management Lead
Global Cyber-Security Team
Shell IT International B.V.



The team

1st

Biggest European
company by revenue
(5th world-wide)

4

IT Hubs across the
world in 3 continents

100+

Cyber Security
professionals, global
in-house team

TB+

Multiple Terabytes
per day on Splunk
ingestion

Agenda



- 1) Effective Vulnerability Management**
How to do a decent job without vulnerability scanners

- 2) Asset & Software inventory**
Enumeration of assets & software, data sources

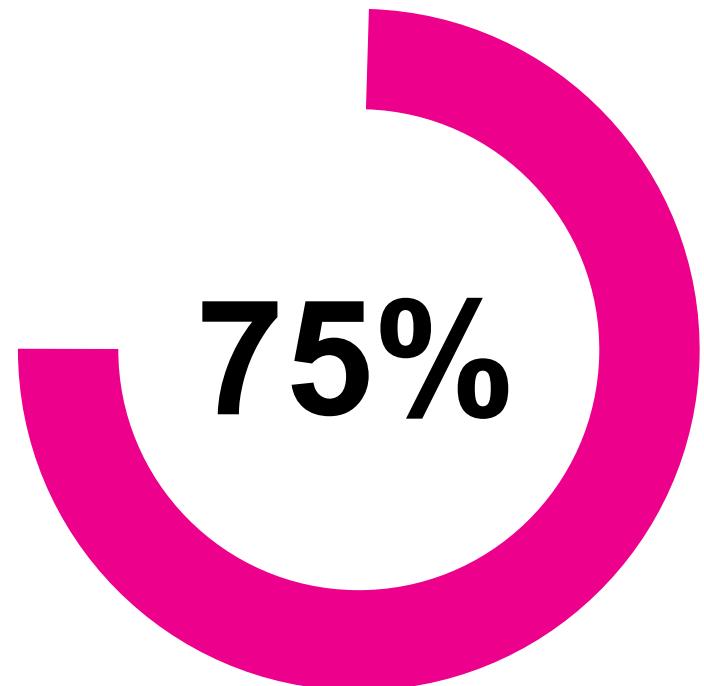
- 3) Data ingestion**
TERM() & PREFIX() ready data

- 4) Data processing**
Cooking the data to make it even more TERM() available

- 5) Data querying**
Use the asset inventory to track vulnerabilities

- 6) Real world examples of usage**

**Asset inventory and Software
inventory value weight on
Vulnerability Management**



It is all about knowing what you have



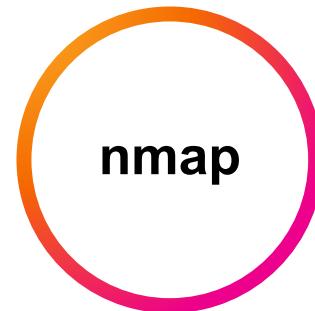
Our experience says up to 75% of cyber-security vulnerability management is just about a solid **Asset inventory & Software inventory**. This means you only need vulnerability scanners for the remaining 25% of vulnerabilities and to automatically follow up on all the universe over time.

Asset inventory sources

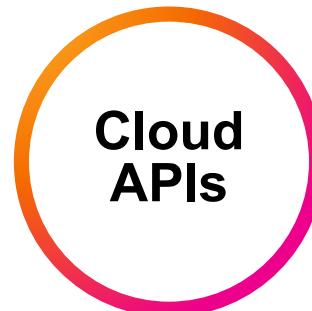
Machine-generated only sources, they all produce a valid Hostname



Traditional
appliance-based
scanners



nmap
'discovery-only'
scripts for
isolated networks
or in case you
don't have
scanners



Enumeration of
VMs, storage,
load balancers,
etc.



EDR agent, AV
agent, scanning
agent, all other
agents you can
read from



Enumeration of
AD domain trust,
later proceed with
enumeration of
Computers per
domain



Always avoid 'human generated' asset
inventory data (aka spreadsheets).

CMDB is not asset inventory

Instead, consider CMDB as ‘a source’ to feed scanners

Avoid following sources to be considered “line-item” generators:

- Excel files.
- CSV files.
- SQL dumps from systems that do not ‘keep state’.
- “Configuration Items” or “CIs”.

Instead, integrate these into Splunk and scan them via scanners (or nmap).

This technique will guarantee your asset inventory “line-items” are always alive assets.

However

Merge problem

//////////
A lot of **|inputlookup**, **|append** or **|join** commands are required. This solution is slow and does not scale on asset inventories with millions of rows and daily refreshes.



The PREFIX() directive

Making all your data |tstats ready

(with just Splunk Enterprise or Splunk Cloud Platform)



The PREFIX() directive game-changer

PREFIX() = Ability to **retrieve** terms from .tsidx!

- Before Splunk 8, we were only capable of **querying** terms with the **TERM()** directive in **|tstats**, but not having them back to the pipeline (retrieve them).
 - In 'OR' conditions, you never knew which term you hit.
 - E.g. TERM(process=powershell.exe) OR TERM(process=cmd.exe)

A screenshot of a Splunk search interface. The search command is:

```
1 | tstats
2     latest(PREFIX(scheduled_time=)) AS latest_schedule_time
3
4     WHERE
5         index=_internal
6         source=/opt/splunk/var/log/splunk/scheduler.log
7         earliest=@d
8         TERM(daniel.ferreira@secret.com)
9         TERM(important-scheduled-report-name)
10
11     BY
12         PREFIX(status=)
13         PREFIX(success=)
14
15     | rename *= AS *
```

The search results show 656 events from 7/23/21 12:00:00.000 AM to 1/19/38 4:14:07.000 AM. The Statistics tab is selected, showing 1 entry. The results table has columns for status and success, with one row labeled "delegated_remote_completion".

Data Models are not needed.

Note there is no 'FROM' clause!

Data source format matters

JSON, CSV, plain text log?

```
1  {  
2      "dnshostname": "WSAMSW89102.addomain.local",  
3      "cn": "WSAMSW89102",  
4      "lastlogontimestamp": "2021-07-15T09:50:07.6943274+02:00",  
5      "operatingsystemversion": "10.0 (18901)",  
6      "operatingsystem": "Windows 10 Enterprise"  
7  }
```



Double quote major breaker? No **PREFIX()** for you! Sorry! 😞

What if we can 'enforce' some key fields to be TERM() and PREFIX() ready in JSON format?

Introducing the ‘dotnotation’ property

Making all JSON data sources **PREFIX()**-ready

```
1 {  
2   "dnshostname": "WSAMSW89102.addomain.local",  
3   "cn": "WSAMSW89102",  
4   "lastlogontimestamp": "2021-07-15T09:50:07.6943274+02:00",  
5   "operatingsystemversion": "10.0 (18901)",  
6   "operatingsystem": "Windows 10 Enterprise",  
7   "dotnotation": [  
8     "cn=WSAMSW89102",  
9     "dnshostname=WSAMSW891025.addomain.local"  
10    ]  
11 }
```

Guaranteeing presence of key fields in **PREFIX()**-ready format.

- Asset name, hostname, FQDN or any other DNS information is always preferable.
 - Asset ID is not mandatory as it will vary per source system.
 - IP address only if no DNS datapoint is seen.
-
- Not meant for ‘all properties’, just ‘key fields’.
 - Trade offs:
 - Some bytes more per event on license.
 - Some bytes more per event on storage.
 - Processing compute power (cheap alternative: server-less).
 - ⚠ Be careful with prefixes clashes.



Get [ConvertTo-DotNotation](#) PowerShell cmdlet from Github.

splunk> .conf21

ConvertTo-SplunkHEC: PowerShell cmdlet

Because we ❤️ languages with pipelines!

```
PS C:\> # Asset inventory from scanner, with dotnotation and converted to HEC payload:  
>> (Import-Csv E:\data\scanner_assets.csv `| Select-Object * -If{$_.dotnotation} | Select-Object asset_id, host_name, ip_address | ConvertTo-DotNotation)) `| ConvertTo-SplunkHEC -IndexValue scanner -SourcetypeValue asset:inventory -SourceValue 'asset_id=(?<source>[\w\-\.\.]+)' -SourceValueIsRegex -HostValue 'ip_address=(?<host>[\w\-\.\.]+)' -HostValueIsRegex -OutputArrayChunks 10000
```

```
1 [  
2 [  
3 {  
4   "event": {  
5     "asset_id": "613135",  
6     "ip_address": "10.45.244.21",  
7     "host_name": "WSAMSW891025.addomain.local",  
8     "os_name": "Windows",  
9     "os_version": "10.0 (18901)",  
10    "os_description": "",  
11    "os_system": "Microsoft Windows",  
12    "last_assessed_for_vulnerabilities": "07/12/2021 18:20:20",  
13    "dotnotation": [  
14      "asset_id=613135",  
15      "host_name=WSAMSW891025.addomain.local",  
16      "ip_address=10.45.244.21"  
17    ],  
18  },  
19  "index": "scanner",  
20  "sourcetype": "asset:inventory",  
21  "source": "613135",  
22  "host": "10.45.244.21",  
23  "time": 1627059093  
24 }]  
]
```

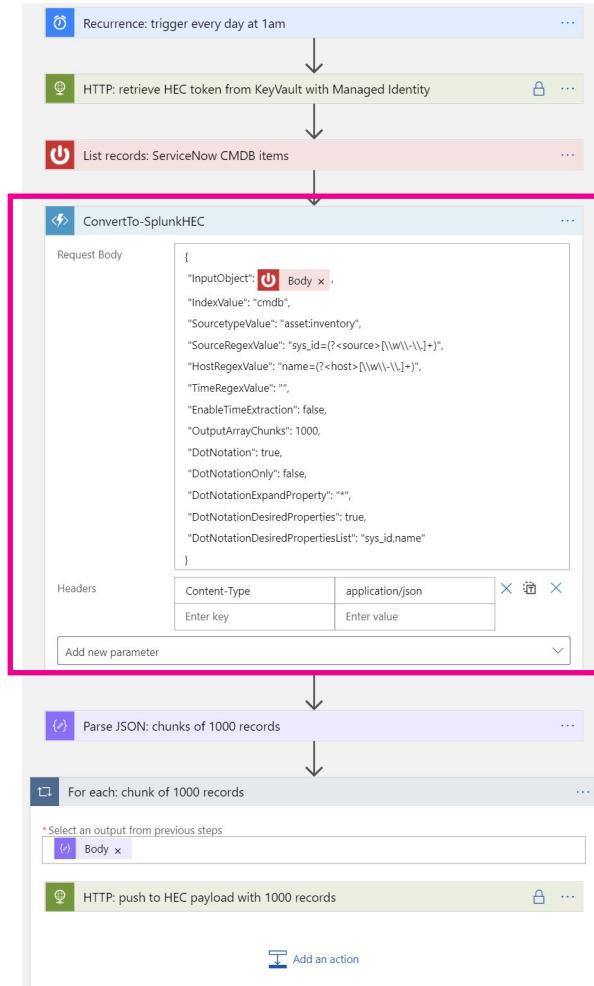
Splunk HEC ready payload
with 'dotnotation' object.



Get [ConvertTo-SplunkHEC](#) PowerShell cmdlet from Github.

Cloud APIs integrations

Leveraging the compute power of PaaS serverless integration services



Cloud-native integration tools help to ‘glue’ 3rd party systems together.

- **Azure Logic Apps with Managed Identities**

- Ideal for batch-based Cloud integration.
- Not only for Azure-based data. Anything that has an API can be plugged in.

- **Azure Functions / AWS Lambda / Google Functions**

- Ideal for streaming-based Cloud integrations.
- Most likely paired with other streaming tools: Event Hubs, Kinesis, etc.

- **Look for affordability:** money is , you can integrate high-volume & high-speed data with few .



Get [ConvertTo-SplunkHEC \(serverless\)](#) PowerShell cmdlet from Github.

Serverless version contains **ConvertTo-DotNotation** cmdlet.

splunk> .conf21

```
1 | tstats
2   fillnull_value=null
3   latest(_time) AS asset_last_seen_epoch
4
5   ``` CMDB: ```
6   latest(PREFIX(sys_id=)) AS cmdb__sys_id
7
8   ``` Scanner: ```
9   latest(PREFIX(asset_id=)) AS scanner_asset_id
10
11 WHERE
12   (index=cmdb OR index=scanner OR index=ad)
13   sourcetype=asset:inventory
14   earliest=-2mon
15
16 BY
17   PREFIX(name=)    `` name:      DNS name coming from CMDB ```
18   PREFIX(cn=)      `` cn:       DNS name coming from Active Directory ```
19   PREFIX(host_name=) `` host_name: DNS name coming from Scanner ```
20   ```` Why not 'host' field? Because it can have different meaning or unexpected value: e.g. UF hostname. ```
21 | rename *= AS *
22
23 ```` Convert "null" to real null() to enable coalesce to work: ```
24 | foreach name,cn,host_name [ | eval <>FIELD>=if(match(<>FIELD>,"^null$"), null(), <>FIELD>) ]
25
26 ```` Coalesce fields across data sources: ```
27 | eval asset_id=coalesce(cmdb__sys_id, scanner_asset_id)
28 | eval asset_name=coalesce(name, cn, host_name)
29
30 ```` Enforce NetBIOS-equivalent DNS name: ```
31 | rex field=asset_name "(?<asset_name>[^.]*"
32
33 ```` Merge: ```
34 | stats max(asset_last_seen_epoch) AS asset_last_seen_epoch, values(asset_id) AS asset_id BY asset_name
```

✓ 3 events (7/22/21 8:37:11.000 PM to 1/19/38 4:14:07.000 AM) No Event Sampling ▾

Events Patterns Statistics (1) Visualization

100 Per Page ▾ Format Preview ▾

asset_name	asset_last_seen_epoch	asset_id
wsamsw89102	1627063915	2de56707148c2e303fbebb645267a284 613135

Alive assets: your asset inventory contains assets seen 'alive' over the last 2 months.

Because you only import 'alive assets as of yesterday', per source, per day; you're naturally embedding 'asset decay' for assets not showing up over the period.

Asset Enrichment Feeds



- 1) CMDB**
 - Business, Sub-business, Service provider, Owner, Deployment, Class, Status, etc.
- 2) Cloud tenants: Azure, AWS, Google**
 - Platform, Owner, Owner by action, Last seen, ID, Local users, Installed software, etc.
- 3) Agents (EDR, AV, sysmon, etc):**
 - OS, Installed software, Installed KBs, Missing KBs, Last reboot, Logged on, etc.
- 4) Active directory:**
 - Last seen, Memberof, Forest, Domain, OUs, custom properties, etc.
- 5) Custom pulls (local scripts, UF scripts, etc):**
 - Installed KBs, Last reboot, Installed software, Local users, etc.
- 6) Custom scans:**
 - HTTP Headers, HTML Title, Redirect URLs, etc.
- 7) Any other reliable asset information...**

However

Merge problem again 😞



How do we merge all of these data repositories? What to use as '**unique key**' field?

Answer: we do not need a single '**unique key**', instead we can have '**many possible unique keys**' 😊



Introducing multi-value: asset_key

All possible ‘unique identifiers’ for an IT asset

😍 Multi-value fields allow you to do many-to-many joins with `|lookup` command.

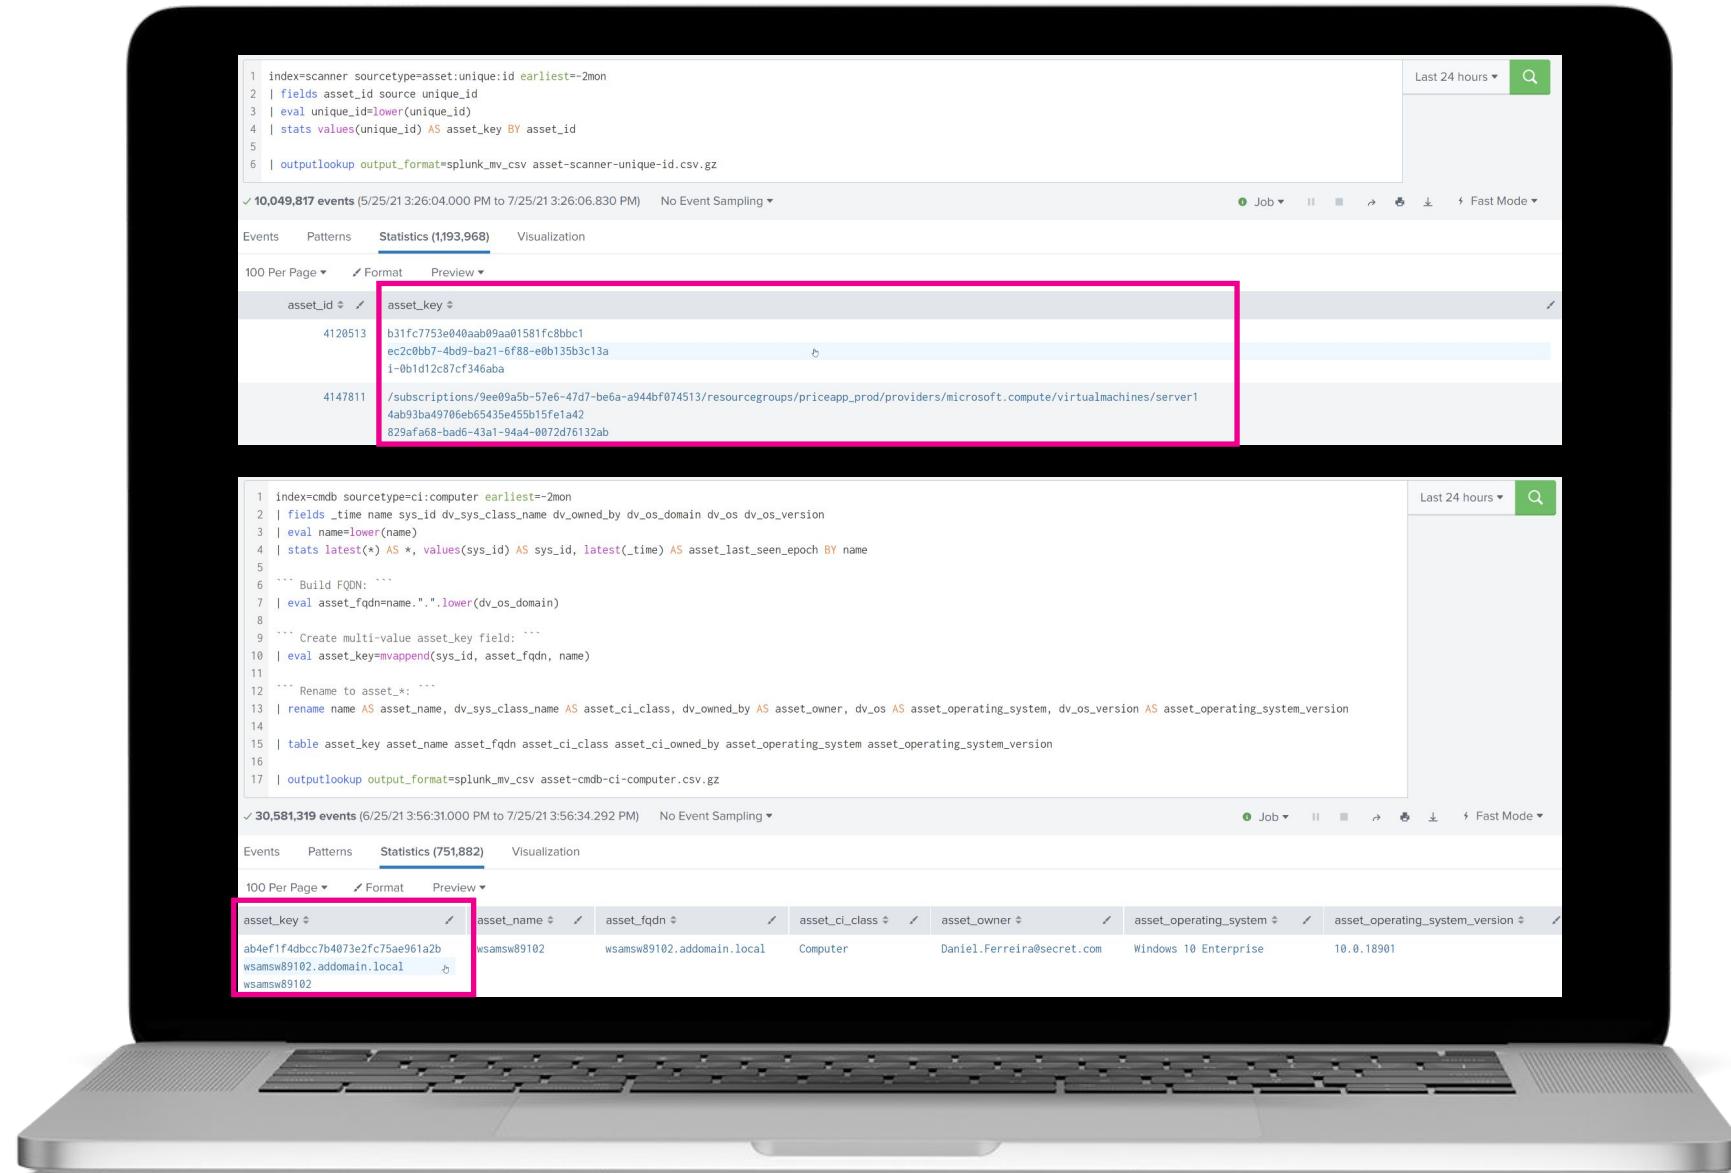
asset_key	
3951416	Scanner ID.
912f79254b8f234bac78c584b3c9e9b2	EDR Agent ID.
10.31.21.77	Private IP address.
52.11.193.95	Public IP address.
ec2amaz-priceapp1	AWS EC2 Instance Name.
ec2-52-11-193-95.ap-southeast-1.compute.amazonaws.com	AWS EC2 Public IP FQDN and short DNS.
ec2-52-11-193-95	AWS EC2 Public IP FQDN and short DNS.
www.pricingapp.domain.com	Domain Name ‘A record’ resolution.
i-0f88f1ee8f16eea41	AWS EC2 Instance ID.
ip-10-31-21-77.ap-southeast-1.compute.internal	AWS EC2 Private IP FQDN.

Integration lookups

Each enrichment source will have its own lookup



Prepare all asset enrichment data in **independent lookup files**, all with its very unique '**asset_key**' multi-value field.



Top Search Result (asset:unique:id):

```

1 index=scanner sourcetype=asset:unique:id earliest=-2mon
2 | fields asset_id source unique_id
3 | eval unique_id=lower(unique_id)
4 | stats values(unique_id) AS asset_key BY asset_id
5
6 | outputlookup output_format=splunk_mv_csv asset-scanner-unique-id.csv.gz

```

✓ 10,049,817 events (5/25/21 3:26:04.000 PM to 7/25/21 3:26:06.830 PM) No Event Sampling ▾

Events Patterns Statistics (1,193,968) Visualization

100 Per Page ▾ Format Preview ▾

asset_id	asset_key
4120513	b31fc7753e040aab09aa01581fc8bbc1 ec2c0bb7-4bd9-ba21-6f88-e0b135b3c13a 1-001d12c87cf346aba
4147811	/subscriptions/9ee09a5b-57e6-47d7-be6a-a944bf074513/resourcegroups/priceapp_prod/providers/microsoft.compute/virtualmachines/server1 4ab93ba49706eb65435c455b15fe1a42 829afaf68-bad6-43a1-94a4-0072d7612ab

Bottom Search Result (cmdb:computer):

```

1 index=cmdb sourcetype=ci:computer earliest=-2mon
2 | fields _time name sys_id dv_sys_class_name dv_owned_by dv_os_domain dv_os_version
3 | eval name=lower(name)
4 | stats latest(*) AS *, values(sys_id) AS sys_id, latest(_time) AS asset_last_seen_epoch BY name
5
6 *** Build FQDN: ***
7 | eval asset_fqdn=name.".lower(dv_os_domain)
8
9 *** Create multi-value asset_key field: ***
10 | eval asset_key=mvappend(sys_id, asset_fqdn, name)
11
12 *** Rename to asset_*: ***
13 | rename name AS asset_name, dv_sys_class_name AS asset_ci_class, dv_owned_by AS asset_owner, dv_os AS asset_operating_system, dv_os_version AS asset_operating_system_version
14
15 | table asset_key asset_name asset_fqdn asset_ci_class asset_ci_owned_by asset_operating_system asset_operating_system_version
16
17 | outputlookup output_format=splunk_mv_csv asset-cmdb-ci-computer.csv.gz

```

✓ 30,581,319 events (6/25/21 3:56:31.000 PM to 7/25/21 3:56:34.292 PM) No Event Sampling ▾

Events Patterns Statistics (751,882) Visualization

100 Per Page ▾ Format Preview ▾

asset_key	asset_name	asset_fqdn	asset_ci_class	asset_owner	asset_operating_system	asset_operating_system_version
ab4ef1f4dbc7b4073e2fc75ae961a2b wsamsrw89102.addomain.local wsamsrw89102	wsamsrw89102	wsamsrw89102.addomain.local	Computer	Daniel.Ferreira@secret.com	Windows 10 Enterprise	10.0.18901

```

25
26     ``` Coalesce fields across data sources: ```
27 | eval asset_id=coalesce(cmdb__sys_id, scanner__asset_id)
28 | eval asset_name=coalesce(name, cn, host_name)
29
30     ``` Enforce NetBIOS-equivalent DNS name: ```
31 | rex field=asset_name "(?<asset_name>[^.]*"
32
33     ``` Merge: ```
34 | stats max(asset_last_seen_epoch) AS asset_last_seen_epoch, values(asset_id) AS asset_id BY asset_name
35
36     ``` asset_key definition: ```
37 | eval asset_key=mvappend(asset_id, asset_name)
38
39     ``` Bring other possible unique-identifiers for the asset when they have agent installed. Coming from scanner, assets with agents installed: ```
40 | lookup asset-scanner-unique-id.csv.gz asset_id AS asset_key OUTPUT asset_key AS temp_scanner_unique_id__asset_key
41 | eval asset_key=mvdedup(mvappend(asset_key, temp_scanner_unique_id__asset_key))
42
43     ``` Bring CMDB enrichment fields: ```
44 | lookup asset-cmdb-ci-computer.csv.gz asset_key OUTPUT asset_key AS temp_cldb__asset_key asset_ci_class asset_owner asset_operating_system asset_operating_system_version
45 | eval asset_key=mvdedup(mvappend(asset_key, temp_cldb__asset_key))
46
47 | table asset_key asset_name a*

```

✓ 3 events (7/22/21 5:11:14.000 PM to 1/19/38 4:14:07.000 AM) No Event Sampling ▾

Events Patterns Statistics (1) Visualization

100 Per Page ▾ Format Preview ▾

asset_key	asset_name	asset_ci_class	asset_id	asset_last_seen_epoch	asset_operating_system	asset_operating_system_version	asset_owner
2de56707148c2e303fbebb645267a284 613135	wsamsw89102	Computer	2de56707148c2e303fbebb645267a284 613135	1627063915	Windows 10 Enterprise	10.0.18901	Daniel.Ferreira@secret.com
wsamsw89102	wsamsw89102.addomain.local						

asset_key: initial definition using all unique identifiers coming from the ‘line-item generation command’ and DNS name.

Many-to-many match: search in each lookup by all possible values seen in ‘asset_key’ field. Then, bring all enrichment fields.

Add more values to ‘asset_key’: each enrichment source may have detected other unique identifiers for the asset, so bring them all into ‘asset_key’ too.

Few months later...

(actually not as happy as these faces)



```

47    ```
48    ### ACTIVE DIRECTORY:
49    ``
50 | lookup asset-active-directory.csv.gz asset_key OUTPUT asset_domain_cn, asset_ad_domain, asset_ad_distinguished_name, asset_ad_last_logon_timestamp, asset_operating_system AS
      temp_ad_asset_os, asset_operating_system_version AS temp_ad_asset_os_version, asset_ad_memberof
51 | eval asset_operating_system=coalesce(temp_ad_asset_os, asset_operating_system)
52 , asset_operating_system_version=coalesce(temp_ad_asset_os_version, asset_operating_system_version)
53 , asset_operating_system_source=if(isnull(asset_domain_cn), asset_operating_system_source, "AD:Computers")    ```: define first field 'asset_operating_system_source'. ````
54 , asset_operating_system_certainty=if(isnull(asset_domain_cn), asset_operating_system_certainty, "1")        ```: define first field 'asset_operating_system_certainty'. ````
55 , asset_is_unlocked=if(isnull(asset_domain_cn), asset_is_unlocked, if(match(lower(asset_ad_distinguished_name),"ou=admin rights") AND match(lower(asset_ad_distinguished_name),"ou=workstations"), "true", asset_is_unlocked))                                ```: detect admin-rights workstations. Define the field first. ````
56 , asset_is_domain_joined=if(isnull(asset_domain_cn),"false","true")
57

58
59    ```
60    ### EDR:
61    ``
62 | lookup asset-edr.csv.gz asset_key OUTPUT asset_key AS temp_edr_asset_key, asset_os_name AS temp_edr_asset_os_name, asset_os_version AS temp_edr_asset_os_version
63 | eval asset_operating_system=coalesce(temp_edr_asset_os_name, asset_operating_system)
64 , asset_operating_system_version=coalesce(temp_edr_asset_os_version, asset_operating_system_version)
65 , asset_operating_system_source=if(isnull(temp_edr_asset_key), asset_operating_system_source, "EDR")
66 , asset_operating_system_certainty=if(isnull(temp_edr_asset_key), asset_operating_system_certainty, "1")
67 , asset_installed_agent=if(isnull(temp_edr_asset_key), asset_installed_agent, mvappend(asset_installed_agent,"EDR-vendor-X"))
68 , asset_key=if(isnull(temp_edr_asset_key), asset_key, mvdedup(mvappend(asset_key,temp_edr_asset_key)))
69
70    ```
71    ### EDR: Last logged on:
72    ``
73 | lookup asset-edr-userlogon.csv.gz asset_key OUTPUT asset_last_logged_on_username AS temp_edr_asset_last_logged_on_username
74 | eval asset_last_logged_on=mvappend(asset_last_logged_on, temp_edr_asset_last_logged_on_username)
75
76    ```
77    ### EDR: Host applied vulnerability kb:
78    ``
79 | lookup asset-edr-installed-kb.csv.gz asset_key OUTPUT asset_installed_vulnerability_kb AS temp_edr_installed_kb
80 | eval asset_installed_vulnerability_kb=mvappend(asset_installed_vulnerability_kb, temp_edr_installed_kb)
81
82    ```
83    ### EDR: Host last reboot::
84    ``
85 | lookup asset-edr-last-reboot.csv.gz asset_key OUTPUT asset_last_reboot_epoch AS temp_edr_asset_last_reboot_epoch
86 | eval asset_last_reboot_epoch=if(isnull(temp_edr_asset_last_reboot_epoch), asset_last_reboot_epoch, temp_edr_asset_last_reboot_epoch)

```

AD Groups this asset belongs to:
remember this for later...

Installed KBs per machine: thereby **the opposite can also be calculated** (assets missing certain KB), in conjunction with '*asset_operating_system*' fields.

```

88     ```
89     ### Azure Virtual Machines:
90
91 | lookup asset-cloud-azure-vms.csv.gz asset_key OUTPUT asset_key AS temp_az__asset_key, asset_platform AS temp_az__asset_platform, asset_platform_description AS temp_az__asset_platform_description, asset_owner_by_action AS
92 |   temp_az__asset_owner_by_action, asset_owner AS temp_az__asset_owner, asset_installed_agent AS temp_az__asset_installed_agent
93 | eval asset_platform=if(isnull(temp_az__asset_key), asset_platform, temp_az__asset_platform)
94 , asset_platform_description=if(isnull(temp_az__asset_key), asset_platform_description, temp_az__asset_platform_description)
95 , asset_owner=mvappend(asset_owner,temp_az__asset_owner)
96 , asset_owner_by_action=if(isnull(temp_az__asset_owner_by_action), asset_owner_by_action, mvdedup(mvappend(asset_owner_by_action,temp_az__asset_owner_by_action))) *** 😊 ***
97 , asset_key=if(isnull(temp_az__asset_key), asset_key, mvdedup(mvappend(asset_key,temp_az__asset_key)))
98 , asset_installed_agent=mvdedup(mvappend(asset_installed_agent, temp_az__asset_installed_agent))
99
100    ```
100  ### AWS Virtual Machines:
101
102 | lookup asset-cloud-aws-vms.csv.gz asset_key OUTPUT asset_key AS temp_aws__asset_key, asset_platform AS temp_aws__asset_platform, asset_platform_description AS temp_aws__asset_platform_description, asset_owner_by_action AS
103 |   temp_aws__asset_owner_by_action, asset_owner AS temp_aws__asset_owner
104 | eval asset_platform=if(isnull(temp_aws__asset_key), asset_platform, temp_aws__asset_platform)
105 , asset_platform_description=if(isnull(temp_aws__asset_key), asset_platform_description, temp_aws__asset_platform_description)
106 , asset_owner=mvdedup(mvappend(asset_owner,temp_aws__asset_owner))
107 , asset_owner_by_action=if(isnull(temp_aws__asset_owner_by_action), asset_owner_by_action, mvdedup(mvappend(asset_owner_by_action,temp_aws__asset_owner_by_action))) *** 😊 ***
108 , asset_key=if(isnull(temp_aws__asset_key), asset_key, mvdedup(mvappend(asset_key,temp_aws__asset_key)))

```

asset_key	asset_name	asset_platform	asset_platform_description
3951416 912f79254b8f234bac78c584b3c9e9b2	ec2amaz-priceapp1	AWS	AWS Account: SD9100039 (Retail Price App - Prod) AWS Instance name: ec2amaz-priceapp1
10.31.21.77			AWS EC2 virtual machine instance id: i-0f88f1ee8f16eea41
52.11.193.95			AWS Public IP: 52.11.193.95
ec2amaz-priceapp1 ec2-52-11-193-95.ap-southeast-1.compute.amazonaws.com ec2-52-11-193-95 www.pricingapp.domain.com i-0f88f1ee8f16eea41 ip-10-31-21-77.ap-southeast-1.compute.internal			

Installed software

Important fields - example

Deduplication: imperfect deduplication of repeated values of software names coming from multiple sources.

Installed software:

- Pulled from custom scripts (e.g. UF PowerShell), EDR, SCCM, Scanner, etc.
 - For asset inventory purposes, keep the 'software name' only **without the version number**.

```

110
111
112     ...
113
114     ### Installed software products:
115     ### Note: the purpose of asset_installed_software is to show installed products, not software versions; so the asset becomes SEARCHABLE by product names.
116     ...
117     ### Installed software products - EDR: ...
118     | lookup software-edr.csv.gz asset_key OUTPUT software_name AS temp_edr__software
119     | eval asset_installed_software=if(isnull(temp_edr__software), asset_installed_software, mvdedup(mvappend(asset_installed_software,temp_edr__software)))
120     ...
121     ### Installed software products - Custom pulling script:
122     | lookup software-uf-pull.csv.gz asset_key OUTPUT software_name AS temp_uf__software
123     | eval asset_installed_software=if(isnull(temp_uf__software), asset_installed_software, mvdedup(mvappend(asset_installed_software,temp_uf__software)))
124     ...
125     ### Installed software products - Scanner (agent and remote authenticated scans):
126     | lookup software-scanner.csv.gz asset_key OUTPUT software_name AS temp_scanner__software
127     | eval asset_installed_software=if(isnull(temp_scanner__software), asset_installed_software, mvdedup(mvappend(asset_installed_software,temp_scanner__software)))
128

```

Owner By Action

Important fields - example

Owner by action:

- Anybody that ‘touches’ a resource within a Cloud Subscription/Account becomes automatically an owner.
 - Need to onboard to Splunk the Audit Logs for all Azure Subscriptions / AWS Accounts.
(you should onboard this regardless)

Keeping difference between owners:
officially registered owners should keep
their own fields, so ‘discovered’ ones
can be identified.

asset_key	asset_platform	asset_platform_description	asset_owner	asset_owner_key	asset_comment
11ab9f3e57a7470db9d729200b538a51 4156479 conq-eh-test 10.201.21.40 conq-eh-test.adddomain.local a1277570-cd85-486f-a7ef-c077c00ac4b1	Azure	Azure tenant: aztenant.onmicrosoft.com Azure subscription id: 8180ab5b-d0c3-41a7-8e91-220ed8d721bb Azure subscription name: Conq-App-TEST Azure resource group: ConQ-EH-TEST-521341 Azure virtual machine: ConQ-EH-TEST	unknown	daniel.ferreira@secret.com	Owner-by-action found: owner(s) 'daniel.ferreira@secret.com' found performing administrative actions on Azure Subscription 'Conq-App-TEST' and Resource Group 'ConQ-EH-TEST-521341'.

What are we going to do with this?

We are going to make all this data ‘searchable’ by writing to a [Summary Index](#)

summary index

Documentation / Splexicon / Summaryindex

* A B C D E F G H I J K L M N O P R S T U V W

noun

Daily scheduled report...

A special [index](#) that stores the [results](#) of a [scheduled report](#), when you enable summary indexing for the report. Summary indexing lets you [run fast searches](#) over large data sets by spreading out the cost of a computationally expensive report over time. To achieve this, the search that populates the summary index runs on a frequent, [recurring basis](#) and [extracts the specific data that you require](#). You can then [run fast and efficient searches](#) against this small subset of data in the summary index.

Daily... **Enrichment fields** brought up by |lookup commands... `asset` macro with natural language text...

Freeze asset inventory data daily

Usage of `|collect` command

```

2148 ...
2149 ...
2150     ### Set _time as of "today at the top of the day" and increase seconds each 100 rows to avoid "sub-second" error:
2151 ...
2152 | streamstats count as row
2153 | eval _time=relative_time(now(),"@d")+(row/100)
2154 ...
2155 ...
2156     ### Convert all Multi-value fields into non-mv by using the " | " delimiter.
2157 ...
2158 | foreach asset* [ eval <<FIELD>>=if(mvcount(<<FIELD>>)>1,mvjoin(<<FIELD>>, " | "),<<FIELD>>) ]
2159 ...
2160 ...
2161     ### Table desired fields:
2162 ...
2163 | table asset_key asset*
2164 ...
2165 ...
2166     ### COLLECT:
2167 ...
2168 | collect index=summary source=asset

```



... a few months later ...

`asset` macro definition:
index=summary source=asset earliest=-01d@d latest=@d

Why the 'earliest' and 'latest'? Because the idea is that Asset Inventory Summary Index gets updated once a day, and all automations required for it will take a while to execute (e.g. AD integration), potentially hours. This timing problem will make ASIA-PAC users unable to search at their morning. So, Asset Inventory in `asset` macro is "updated as of yesterday" by default (unless you resolve it manually and move it to earliest=@d). You should also create a KV lookup containing updated data as of today, which then can feed Splunk ES —if you have it—and allow users to `|lookup` it in their other searches.

Summary index results

Stash format, automatic field extraction

💡 TERM() & PREFIX() ready fields:
Summary Index data is in 'stash format', which will
make your asset inventory data **|tstats ready**.

The screenshot shows a Splunk search interface with the following details:

- Search Bar:** index=summary source=asset
- Time Range:** Last 24 hours
- Event Count:** 1 event (7/25/21 12:28:21.000 PM to 7/26/21 12:28:21.000 PM)
- Sampling:** No Event Sampling
- Job Controls:** Job ▾, II, ■, →, ↵, ↓, Verbose Mode ▾
- Event View:** Events (1) | Patterns | Statistics | Visualization
- Format Timeline:** Format Timeline ▾, Zoom Out, Zoom to Selection, Deselect, 1 hour per column
- Event Details:** Raw ▾, Format, 50 Per Page ▾
- Selected Fields:** asset_exposure 1, asset_key 1, asset_name 1, asset_operating_system 1, asset_operating_system_version 1, asset_owner 1
- Event Data:** 07/26/2021 00:00:00 +0200, info_min_time=1627208888.000, info_max_time=1627295288.000, info_search_time=1627295288.131, asset_key="ab4ef1f4dbcc7b4073e2fc75ae961a2b | wsamsw89102.addomain.local | wsamsw89102". asset_operating_system_version="10.0.18901", asset_operating_system="Windows 10 Enterprise", asset_ci_class=Computer, asset_exposure=intranet, asset_fqdn="wsamsw89102.addomain.local", asset_name=wsamsw89102, asset_owner="daniel.ferreira@secret.com"

“TERM() is used in less than 1% of all customer searches executed on Splunk Cloud.”

Richard Morgan
Principal Architect
Splunk



Summary index queries

|tstats all your data with TERM() and PREFIX()

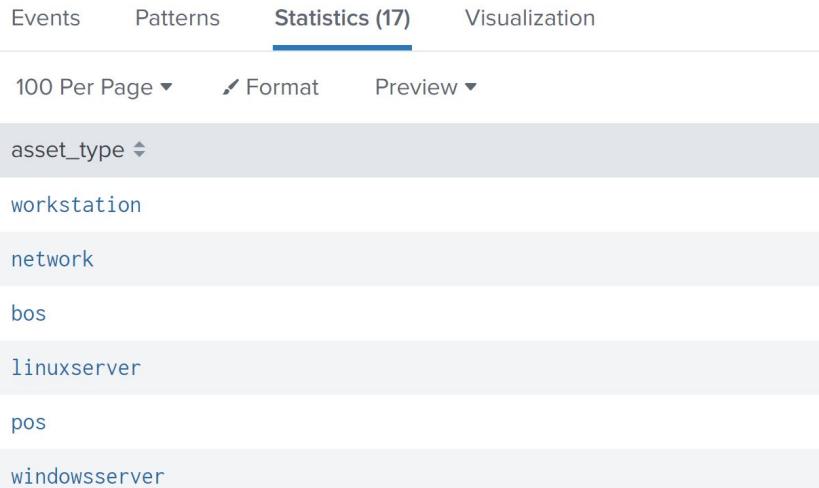
```

1 | tstats count
2   WHERE
3     index=summary
4     source=asset
5     earliest=-1d@d
6     latest=@d
7     BY
8       PREFIX(asset_type=)
9   | rename *= AS *

```

✓ 1,312,195 events (7/25/21 12:00:00.000 AM to 7/26/21 12:00:00.000 AM)

No Event Sampling ▾



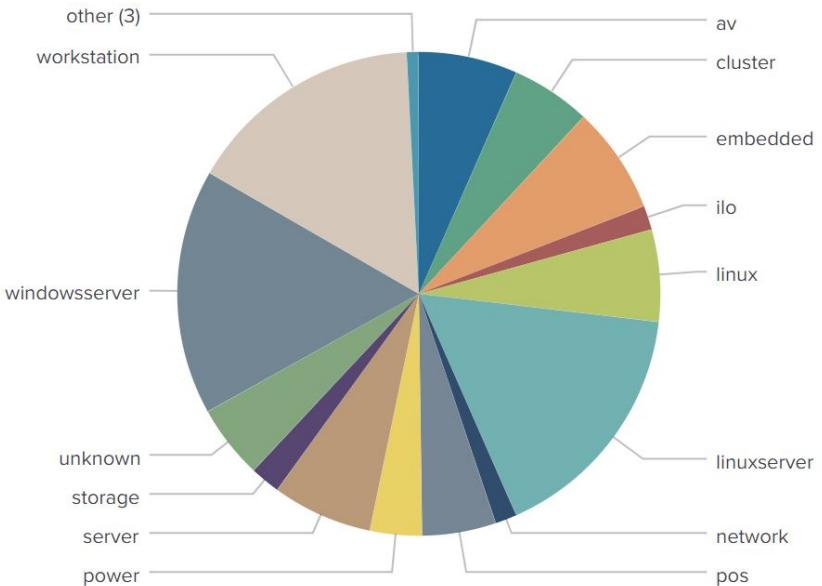
```

1 | tstats count
2   WHERE
3     index=summary
4     source=asset
5     earliest=-30d@d
6     latest=@d
7     TERM(asset_type=windowsserver)
8     BY
9       _time span=1d

```

✓ 3,112,626 events (7/16/21 12:00:00.000 AM to 7/26/21 12:00:00.000 AM)

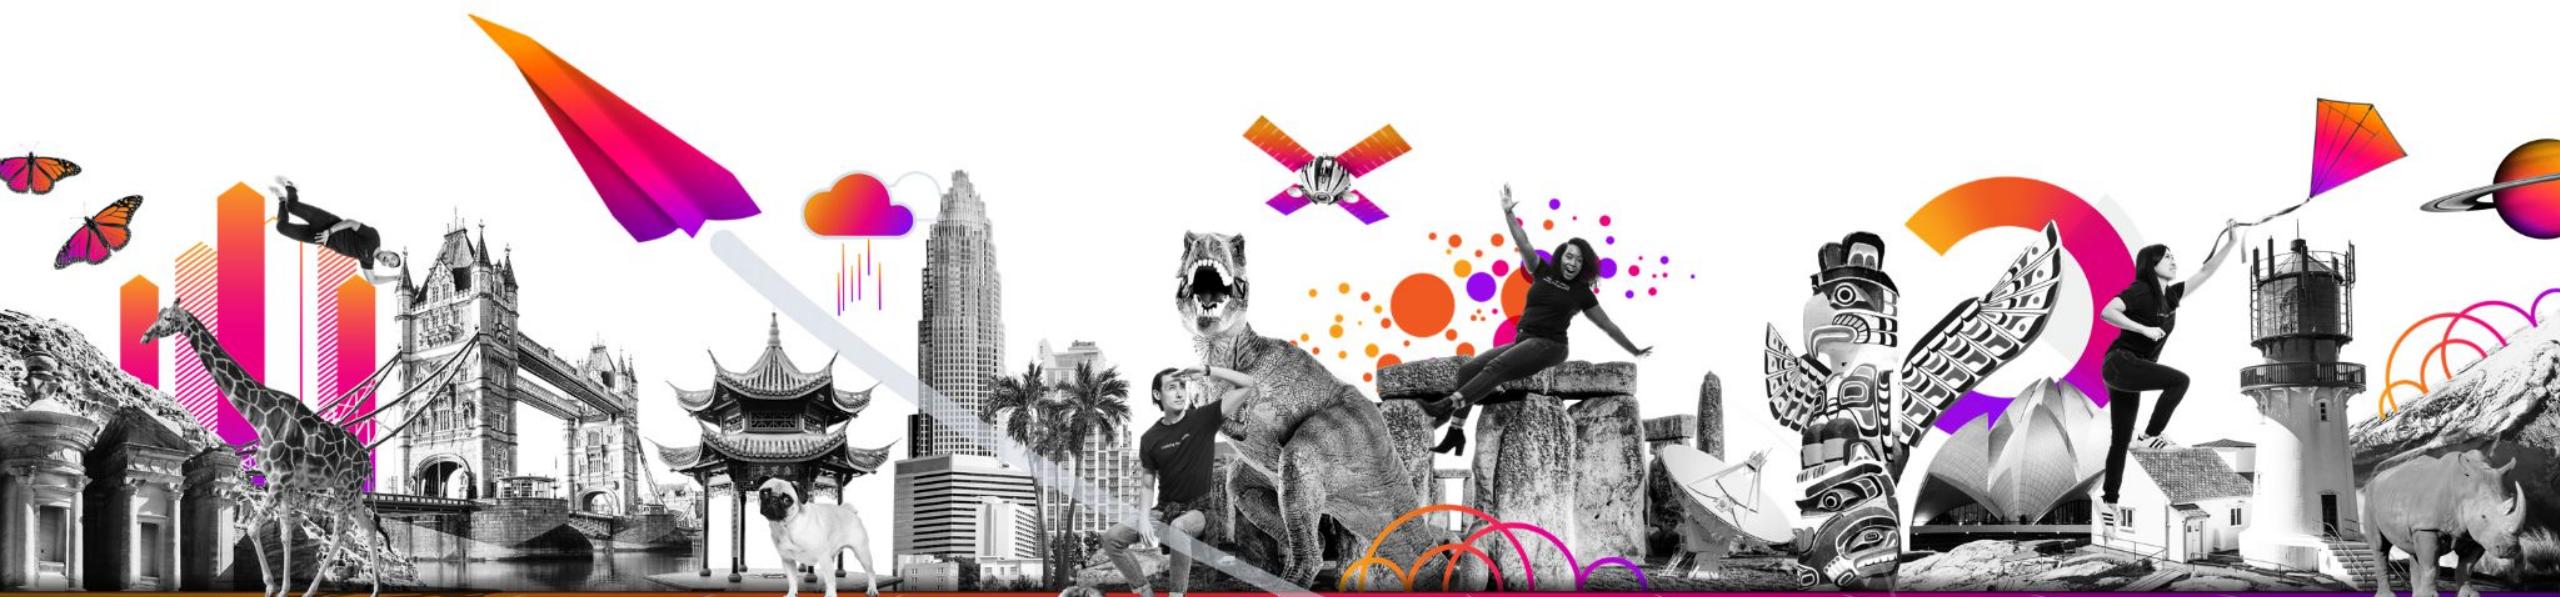
This search has completed and has returned 30 results by scanning 3,112,626 events in 1.839 seconds



313,104 ↘
-156

Few real world examples

(actually happy as these faces)



Asset inventory search examples

Use-case: **breaching report** on an **internet facing** application

```
1 `asset` pricingapp.domain.com
```

```
1 `asset` 52.11.193.95
```

```
1 `asset` retail app
```

```
1 `asset` aws virtual machine public ip prod*
```

```
1 `asset` i-0f88f1ee8f16eea41 ec2
```

```
1 `asset` 921f79254b8f234bac78c584b3c9e9b2
```

```
1 `asset` aws retail machine
```

```
1 `asset` 10.31.21.77
```

```
1 `asset` sd91*
```

asset_key	asset_name	asset_platform	asset_platform_description
3951416	ec2amaz-priceapp1	AWS	AWS Account: SD9100039 (Retail Price App - Prod)
912f79254b8f234bac78c584b3c9e9b2			AWS Instance name: ec2amaz-priceapp1
10.31.21.77			AWS EC2 virtual machine instance id: i-0f88f1ee8f16eea41
52.11.193.95			AWS Public IP: 52.11.193.95
ec2amaz-priceapp1			
ec2-52-11-193-95.ap-southeast-1.compute.amazonaws.com			
ec2-52-11-193-95			
www.pricingapp.domain.com			
i-0f88f1ee8f16eea41			
ip-10-31-21-77.ap-southeast-1.compute.internal			

Asset inventory search examples

Use-case: hunting for 'solarwinds' software

Search by software: assets become 'searchable' by their installed software.

```
1 `asset` solarwinds OR (solar winds)
2 | table asset_key asset_installed_software asset_owner_key
```

Real search: this is a real search we ran when Solarwinds supply-chain attack came up. 😱

The screenshot shows a Splunk search interface with three search terms in the search bar: 'asset', 'solarwinds OR (solar winds)', and 'table asset_key asset_installed_software asset_owner_key'. The 'asset_installed_software' term is highlighted with a pink box. Below the search bar, the results table has columns for 'asset_key', 'asset_installed_software', and 'asset_owner_key'. The 'asset_installed_software' column contains a long list of SolarWinds product names, many of which are also highlighted with pink boxes. The list includes: solarwinds network automation manager, solarwinds database performance analyzer integration module, solarwinds orion quality of experience, solarwinds orion log viewer, solarwinds log analyzer, solarwinds user device tracker, solarwinds virtual infrastructure monitor, solarwinds voip and network quality manager, solarwinds server & application monitor, solarwinds orion network configuration manager, solarwinds orion netflow traffic analyzer, solarwinds netpath, solarwinds administration service, solarwinds active diagnostics, solarwinds ip address manager, solarwinds orion core services, solarwinds information service, solarwinds agent, solarwinds orion network performance monitor, solarwinds rabbitmq server, nmap npcap, solarwinds orion syslog/traps, solarwinds scp server, oracle open java development kit, microsoft monitoring agent, microsoft windows server, microsoft office access database engine, solarwinds orion network atlas, microsoft azure active directory authentication library for sql server, winpcap team winpcap, solarwinds tftp server, microsoft internet explorer, microsoft sql server management objects, microsoft sql server system clr types, microsoft windows media player, solarwinds job engine, and microsoft sql server native client.

Asset inventory search examples

Use-case: CVE-2020-1350 (SIGRed) **DNS TCP RCE vulnerability** (CVSS: 10)

```
1 `asset` asset_exposure=internet tcp/53 windows server  
2 | table asset_key a*ports a*operating_system a*owner_key
```

asset_key	asset_available_ports	asset_operating_system	asset_owner_key
5250671 ddns01.domain.com	tcp/53 tcp/80 tcp/443 tcp/853 udp/53	Windows Server 2019	alice@secret.com

Asset inventory search examples

Use-case: compromised account, **where else is this account seen?** (reverse search)
(e.g. pass-the-hash, ransomware/lateral movement attempt)

```
1 `asset` daniel.ferreira
2 | table asset_key asset_owner_key asset_last_logged_on
3 | `recover-mv`
```

asset_key	asset_owner_key	asset_last_logged_on
010b1dcc3877404a8bccf37e948d4aba	daniel.ferreira@secret.com	bloodhound
5840891		prv.daniel.ferreira
azrsrv1415		
10.21.11.18		
azrsrv1415.addomain.local		

Asset inventory search examples

Use-case: estate hygiene, **asset at risk, missing EDR tool**

```
1 `asset` domain controllers missing-EDR  
2 | table asset_key asset_ad_distinguished_name asset_is_edr*
```

```
1 `asset` TERM(asset_exposure=extranet) TERM(asset_is_business_critical=true) TERM(asset_is_edr_in_scope=true) TERM(asset_is_edr_seen=false)  
2 | table asset_key asset_owner_key asset_last_seen asset_installed_agent  
3 | `recover-mv`
```

TERM() is your friend: note that if you don't use the TERM() directive in this search, you'll effectively search for 'extranet', 'true' and 'false' LISPY terms only.

Not using **TERM()**:

```
[ AND extranet false index::summary source::asset true ]
```

Using **TERM()**:

```
[ AND asset_exposure=extranet asset_is_business_critical=true asset_is_edr_in_scope=true  
asset_is_edr_seen=false index::summary source::asset ]
```

Without **TERM()**, search will return same result but will be **considerable slower!** 😱

Takeaways

- Main generation command for asset inventory must be `|tstats`: allows to handle millions of records.
 - Starting 8.1: command limitation of ~1500 commands per query. Use comma-separated `|eval` for creating fields.
- Data quality is key: ensure a perfect '**host**' field across all data sources, otherwise leverage **ConvertTo-DotNotation** cmdlet (or other method) to add 'key' fields as **PREFIX()**-ready fields.
- Push assets from data sources 'seen/scanned' as of 'yesterday':
 - E.g. Today at 00:00, pull AD Computers with a 'lastlogontimestamp' updated as of the last -24h. AD will refresh this every 14 days.
 - E.g. Today at 00:00, pull assets scanned/discovered by the scanners as of the last -24h. Ensure you scan whole estate every week.
- Before merging different asset inventories, convert all IP addresses to hostname.
- Rely as much as possible on 'machine-generated data'.
 - EDR, AVs, any other agent you may have. Same approach: no 'full-dump' load, instead an incremental one based on 'last seen'. Assets will decay themselves.
- Don't rely on CMDB data unless it comes with a reliable machine-generated 'last seen alive' property.
- Enrich your asset inventory as much as possible. Each field you add will become searchable! 😊
 - Splunk's best value is the ability to **NOT** need `field=value` searches. Make your field values able to be searched this way.
- Write your asset inventory daily both to Summary Index and KV lookup. Consider to feed Splunk ES with this.
 - But always avoid querying your asset inventory with `|inputlookup`!
- Think very well the values of certain fields: make sure you add relevant 'keywords' to them.
 - E.g. "Azure VM", "AWS load balancer", "Detected Internet facing web server hosted in Azure, port tcp/443 exposed" → `asset` web server azure internet exposed.
- The values of your fields will become '**TERM()** and **PREFIX()**-able' if you don't use minor breakers.
 - Very useful for `*_is_*` fields!



Please upvote these on Splunk Ideas:

- Add "OUTPUTMERGE" option to `|lookup`.
- `|collect` command to not double quote fields with just minor breakers.



Thank You

Please provide feedback via the
SESSION SURVEY

